

Compiler Design - Notes Week 6

Ruben Schenk, ruben.schenk@inf.ethz.ch

October 26, 2021

8. LL & LR Parsing

8.1 LL(1) Grammars

One problem with grammars is that not all grammars can be parsed “top-down” with a *single lookahead*. *Top-down* means that we start from the start symbol, i.e. the root of the parse tree, and go down.

LL(1) means:

- Left-to-right scanning
- Left-most derivation
- 1 lookahead symbol

8.1.1 Making a Grammar LL(1)

The main problem is that we can’t decide which S production to apply until we see the symbol after the first expression. The solution is to *left-factor* the grammar. There is a common S prefix for each choice, so add a new non-terminal S' at the decision point:

This means that we transform our example grammar

$$S \rightarrow E + S \mid EE \rightarrow \text{number} \mid (S)$$

to the following “left-factor” grammar:

$$S \rightarrow ES'S' \rightarrow \epsilon S' \rightarrow +SE \rightarrow \text{number} \mid (S)$$

However, we also need to *eliminate left-recursion* somehow. In general, this is done by rewriting the following left-recursive rule

$$S \rightarrow S\alpha_1 \mid \dots \mid S\alpha_n \mid \beta_1 \mid \dots \mid \beta_m$$

to a rule of the form:

$$S \rightarrow \beta_1 S' \mid \dots \mid \beta_m S' S' \rightarrow \alpha_1 S' \mid \dots \mid \alpha_n S' \mid \epsilon$$

In our running example, this would mean to rewrite

$$S \rightarrow S + E \mid EE \rightarrow \text{number} \mid (S)$$

to the following left-recursion-eliminating grammar:

$$S \rightarrow ES'S' \rightarrow +ES' \mid \epsilon E \rightarrow \text{number} \mid (S)$$

8.1.2 Predictive Parsing

Given an LL(1) grammar:

- For a given non-terminal, the lookahead symbol uniquely determines the production to apply
- The parsing is driven by a predictive parsing table
- It is convenient to add a special *end-of-file token* \$ and a start symbol T that requires \$

Example: Let us look at the following LL(1) grammar:

$$T \rightarrow S\$S \rightarrow ES'S' \rightarrow \epsilon S' \rightarrow +SE \rightarrow \text{number} \mid (S)$$

We then propose the following **predictive parsing table**:

	number	+	()	\$ (EOF)
T	$\rightarrow S$				
S	$\rightarrow ES'$				
S'		$\rightarrow +S$			
E	$\rightarrow \text{number}$		$\rightarrow (S)$		

8.1.3 Construction of Parse Table

How do we construct the parse table? We examine two possible cases by considering the following production:
 $A \rightarrow \gamma$

Case 1 Construct the set of all input tokens that may appear *first* in strings that can be derived from γ . Then add the production $\rightarrow \gamma$ to the entry (A, token) for each such token.

Case 2 If γ can derive ϵ , then we construct the set of all input tokens that may *follow* the nonterminal A in the grammar. We then add the production $\rightarrow \gamma$ to the entry (A, token) for each such token.

Note: If there are two different productions for a given entry, then the grammar is not LL(1).

Example:

- $\text{First}(T) = \text{First}(S)$
- $\text{First}(S) = \text{First}(E)$
- $\text{First}(S') = \{ +, \epsilon \}$
- $\text{First}(E) = \{ \text{number}, '(' \}$

- $\text{Follow}(S') = \text{Follow}(S)$
- $\text{Follow}(S) = \{ \$, ')' \} \cup \text{Follow}(S')$

$T \mapsto S\$$
 $S \mapsto ES'$
 $S' \mapsto \epsilon$
 $S' \mapsto +S$
 $E \mapsto \text{number} \mid (S)$

Note: we want the *least* solution to this system of set equations... a *fixpoint* computation. More on these later in the course.

	number	+	()	\$ (EOF)
T	$\mapsto S\$$		$\mapsto S\$$		
S	$\mapsto ES'$		$\mapsto ES'$		
S'		$\mapsto +S$		$\mapsto \epsilon$	$\mapsto \epsilon$
E	$\mapsto \text{number}$		$\mapsto (S)$		

8.1.4 Converting The Parsing Table to Code

When we want to convert a parsing table to code, we proceed as follows:

- Define N mutually recursive functions, one for each non-terminal A : **parse_A**
- **parse_A** is of type **unit** \rightarrow **ast** if $\$A\$$ is not an auxiliary non-terminal
- Otherwise, **parse_A** takes additional ast's as inputs, one for each non-terminal in the “factored” prefix

Then, each function **parse_A**:

- “Peeks” at the lookahead token
- Follows the production rule in the corresponding entry
- Consumes the terminal tokens from the input stream
- Calls `parse_X` to create the sub-tree for the non-terminal X
- If the rule ends in an auxiliary non-terminal, it is called with the appropriate ast’s
- Otherwise, this function builds the ast tree itself and returns it

8.1.5 LL(1) Summary

LL(1) is top-down based parsing that finds the left-most derivation. The process proceeds with the following steps:

1. Language grammar \Rightarrow
2. LL(1) grammar \Rightarrow
3. Prediction table \Rightarrow
4. Recursive-descent parser

8.2 LR Grammars

8.2.1 Bottom-up Parsing

LR(k) parser are *bottom-up parser*:

- Left-to-right scanning
- Right-most derivation
- k lookahead symbols

LR grammars are *more expressive* than LL grammars:

- They can handle left-recursive (and right-recursive) grammars (i.e. virtually all programming languages)
- They make it easier to express programming language syntax (e.g. no left factoring)

The most common technique are **Shift-Reduce parsers**:

- Work bottom up instead of top down
- Construct the right-most derivation of a program in the grammar
- Better error detection/recovery

8.2.2 Shift/Reduce Parsing

In shift/reduce parsing, the parser has a **parser state** described as follows:

- Stack of terminals and non-terminals
- Unconsumed input is a string of terminals
- The current derivation step is **stack + input**

Parsing is a sequence of **shif** and **reduce** operations:

- Shift: Move lookahead token to the stack
- Reduce: Replace symbols γ at the top of the stack with a non-terminal X s.t. $X \rightarrow \gamma$ is a production, i.e. **pop gamma, push X**

Example: We consider our previous example

$$S \rightarrow S + E \mid EE \rightarrow \text{number} \mid (S)$$

Stack	Input	Action
	(1 + 2 + (3 + 4)) + 5	shift (
(1 + 2 + (3 + 4)) + 5	shift 1
(1	+ 2 + (3 + 4)) + 5	reduce: $E \mapsto \text{number}$
(E	+ 2 + (3 + 4)) + 5	reduce: $S \mapsto E$
(S	+ 2 + (3 + 4)) + 5	shift +
(S +	2 + (3 + 4)) + 5	shift 2
(S + 2	+ (3 + 4)) + 5	reduce: $E \mapsto \text{number}$

8.3 LR(0) Grammars

8.3.1 LR Parser States

Our goal it is to know *what set of reductions are legal* at any given point. The idea to solve this problem is to summarize all possible stack prefixes α as a finite parser state:

- The parser state is computed by a DFA that reads the stack σ
- Accept states of the DFA correspond to unique reductions that apply

8.3.2 Example LR(0) Grammar: Tuples

The following grammar is an example grammar for non-empty tuples and identifiers:

$$S \rightarrow (L) \mid \text{id}L \rightarrow S \mid L, S$$

Now, if we apply parsing as a sequence of shift and reduce operations, we end up with the following parse operation:

- **Shift:** Move look-ahead token to the stack

Stack	Input	Action
	(x, (y, z), w)	shift (
(x, (y, z), w)	shift x

- **Reduce:** Replace symbols γ at top of stack with nonterminal X s.t.
 $X \mapsto \gamma$ is a production, i.e., pop γ , push X

Stack	Input	Action
(x	, (y, z), w)	reduce $S \mapsto \text{id}$
(S	, (y, z), w)	reduce $L \mapsto S$

8.3.3 Action Selection Problem

Given a stack σ and a lookahead symbol b , should the parser either:

- Shift b onto the stack (new stack is σb), or
- Reduce a production $X \rightarrow \gamma$, assuming that $\sigma = \alpha\gamma$ (new stack is αX) ?

The main idea to solve this problem is to decide based on a prefix α of the stack plus the lookahead. The prefix α is different for different possible reductions since in productions $X \rightarrow \gamma$ and $Y \rightarrow \beta$, γ and β might have different lengths.

8.3.4 LR(0) States

An **LR(0) state** consists of items to track progress on possible upcoming reductions. An **LR(0) item** is a production with an extra separator $.$ in the RHS. Example items could be: $S \rightarrow \cdot(L)$ or $S \rightarrow (\cdot L)$ or $L \rightarrow S$.

The intuition for the meaning of the dot is:

- Stuff before the $.$ is already on the stack
- Stuff after the $.$ is what might be seen next

8.3.5 Constructing The DFA

We will consider the following grammar:

$$S' \rightarrow S\$S \rightarrow (L) \mid \text{id}L \rightarrow S \mid L, S$$

The first step when creating the DFA is to add a new production $S' \rightarrow S$ to the grammar. The *start state* of the DFA is the empty stack, so it contains the item $S' \rightarrow \cdot S$. We then proceed to add the **closure of the state** to our DFA:

- Add items for all productions whose LHS non-terminal occurs in an item in the state just after the $.$ (e.g. S in $S' \rightarrow \cdot S$)

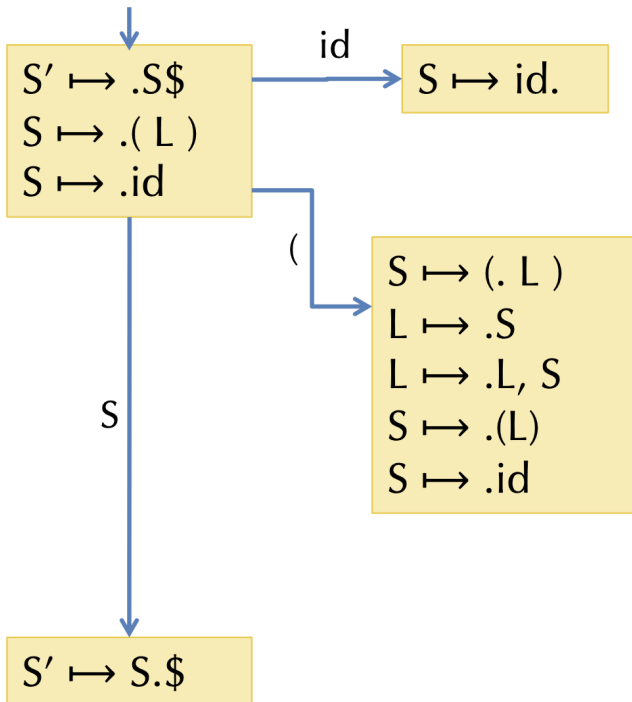
- The added items have the \cdot located at the beginning
- Note that newly added items may cause yet more items to be added to the state, we keep iterating until a *fixed point* is reached

Example:

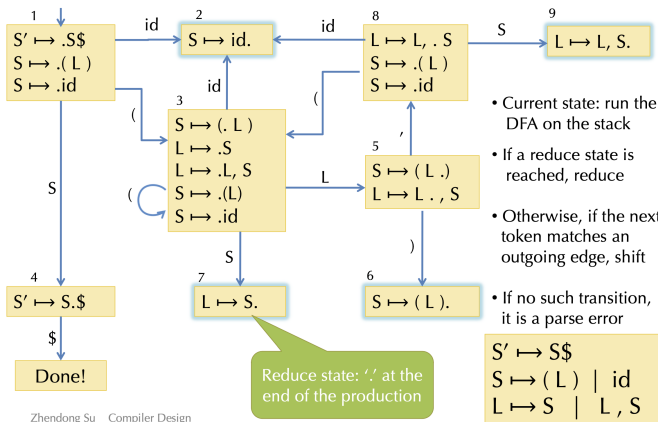
$$\text{Closure}(\{S' \rightarrow .S\}) = \{S' \rightarrow .S, S \rightarrow \cdot(L), S \rightarrow \cdot id\}$$

Next we need to add the *transitions*:

1. First, we see what terminals and non-terminals can appear after the \cdot in the source state.
2. The target state initially includes all items from the source state that have the edge-label symbol after the \cdot , but we advance the \cdot to simulate shifting the item onto the stack.
3. Finally, for each new state, we again take the closure of it.



By continuing the above approach, we will reach the following **full DFA** for our example:



8.3.6 Using The DFA

To use our DFA now, we run the parser stack σ through the DFA. The resulting state tells us what productions may be reduced next:

- If not in a reduce state, we shift the next symbol and transition w.r.t. the DFA
- If in a reduce state, $X \rightarrow \gamma$ with stack $\alpha\gamma$, we **pop** γ and **push** X

Optimization: There is no need to rerun the DFA from the beginning at each step. We might simply store the state with each symbol on the stack, e.g. ${}_1(3)_3L_5)_6$. Then:

- On a reduction $X \rightarrow \gamma$, we **pop** the stack to reveal the state too, e.g. from stack ${}_1(3)_3L_5)_6$ we reduce $S \rightarrow (L)$ to reach stack ${}_1(3)$
- Next, we push the reduction symbol, e.g. to reach the stack ${}_1(3)S$
- Then we take just one step in the DFA to find the next state ${}_1(3)S_7$

8.3.7 Implementing The Parsing Table

We represent the DFA as a table of shape **state** * (**terminals** + **nonterminals**).

Entries for the **action table** specify two kinds of actions:

- Shift and go to state n
- Reduce using the reduction $X \rightarrow \gamma$: First, **pop** γ of the stack to reveal the state, second, look up X in the **goto table** and go to that state

Example:

	()	id	,	\$	S	L
1	s3		s2			g4	
2	S \mapsto id	S \mapsto id	S \mapsto id	S \mapsto id	S \mapsto id		
3	s3		s2			g7	g5
4					DONE		
5		s6		s8			
6	S \mapsto (L)	S \mapsto (L)	S \mapsto (L)	S \mapsto (L)	S \mapsto (L)		
7	L \mapsto S	L \mapsto S	L \mapsto S	L \mapsto S	L \mapsto S		
8	s3		s2			g9	
9	L \mapsto L,S	L \mapsto L,S	L \mapsto L,S	L \mapsto L,S	L \mapsto L,S		

sx = shift and go to state x
 gx = go to state x

8.3.8 LR(0) Limitations

An LR(0) machine only works if states with reduce actions have a *single* reduce action. In such states, the machine always reduces, ignoring lookahead.

With more complex grammars, the DFA construction will yield states with *shift/reduce* and *reduce/reduce* problems:

OK

shift/reduce

reduce/reduce

S \mapsto (L).

S \mapsto (L).
L \mapsto .L , S

S \mapsto L ,S.
S \mapsto ,S.

8.4 LR(1) Parsing

The algorithm for **LR(1) parsing** is similar to LR(0) DFA construction:

- LR(1) state is the set of all LR(1) items
- An LR(1) item is an LR(0) item plus a set of lookahead symbols, i.e. $A \rightarrow \alpha.\beta, \mathcal{L}$

However, the **LR(1) closure** is a little more complex:

1. We first form the set of items just as we did in the LR(0) algorithm
2. Whenever a new item $C \rightarrow \alpha.\gamma$ is added, because the item $A \rightarrow \beta.C\delta, \mathcal{L}$ is already in the set, we need to compute its lookahead set \mathcal{M} :
 1. The lookahead set \mathcal{M} includes $\text{FIRST}(\delta)$, i.e. the set of terminals that may start strings derived from δ

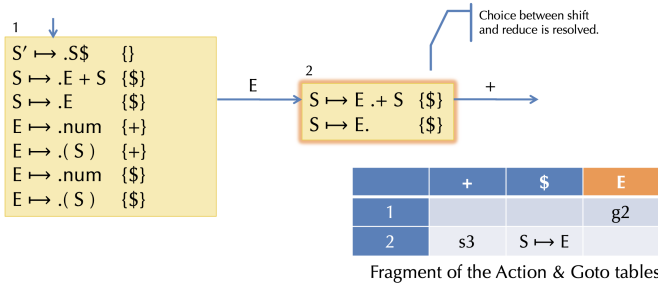
2. If δ is or can derive ϵ , then the lookahead \mathcal{M} also contains \mathcal{L}

8.4.1 Example Closure in LR(1)

$S' \mapsto S\$$
 $S \mapsto E + S \mid E$
 $E \mapsto \text{number} \mid (S)$

- Start item: $S' \mapsto .S\$$, $\{\}$
- Since S is to the right of a '.', add
 $S \mapsto .E + S$, $\{\$ \}$ Note: $\{\$ \}$ is FIRST($\$$)
 $S \mapsto .E$, $\{\$ \}$
- Need to keep closing, since E appears to the right of a '.' in '.E + S'
 $E \mapsto .\text{number}$, $\{+\}$ Note: + added for reason 1
 $E \mapsto .(\text{S})$, $\{+\}$ FIRST(+ S) = $\{+\}$
- Because E also appears to the right of '.' in '.E' we get:
 $E \mapsto .\text{number}$, $\{\$ \}$ Note: \$ added for reason 2
 $E \mapsto .(\text{S})$, $\{\$ \}$ δ is ϵ
- All items are distinct, so we're done

8.4.2 Using The DFA



The behavior is determined if:

- There is no overlap among the lookahead sets for each reduce item, and
- None of the lookahead symbols appear to the right a .

8.4.3 LR Variant: LALR(1)

Consider for example the following two LR(1) states:

$$S_1 : \{[X \rightarrow \alpha., a], [Y \rightarrow \beta., c]\} S_2 : \{[X \rightarrow \alpha., b], [Y \rightarrow \beta., d]\}$$

They have the same core and can therefore be *merged*. The merged state contains:

$$\{[X \rightarrow \alpha., a/b], [Y \rightarrow \beta., c/d]\}$$

These are so-called **LALR(1)** states. Typically, there are 10 times fewer LALR(1) states than LR(1). However, LALR(1) may introduce new reduce/reduce conflicts (but not new shift/reduce conflicts).